

Introduction to OpenCV

Install

- Install Numpy, opencv, matplotlib
 - `conda install numpy matplotlib`
 - `conda install -c menpo opencv`
 - `pip install numpy opencv-python==3.4.2.16 opencv-contrib-python==3.4.2.16 matplotlib`
- Install jupyter notebook
 - `conda install -c conda-forge jupyterlab`
 - `pip install jupyterlab`
- Open jupyter notebook in you desired folder
 - `jupyter lab`

Getting Started

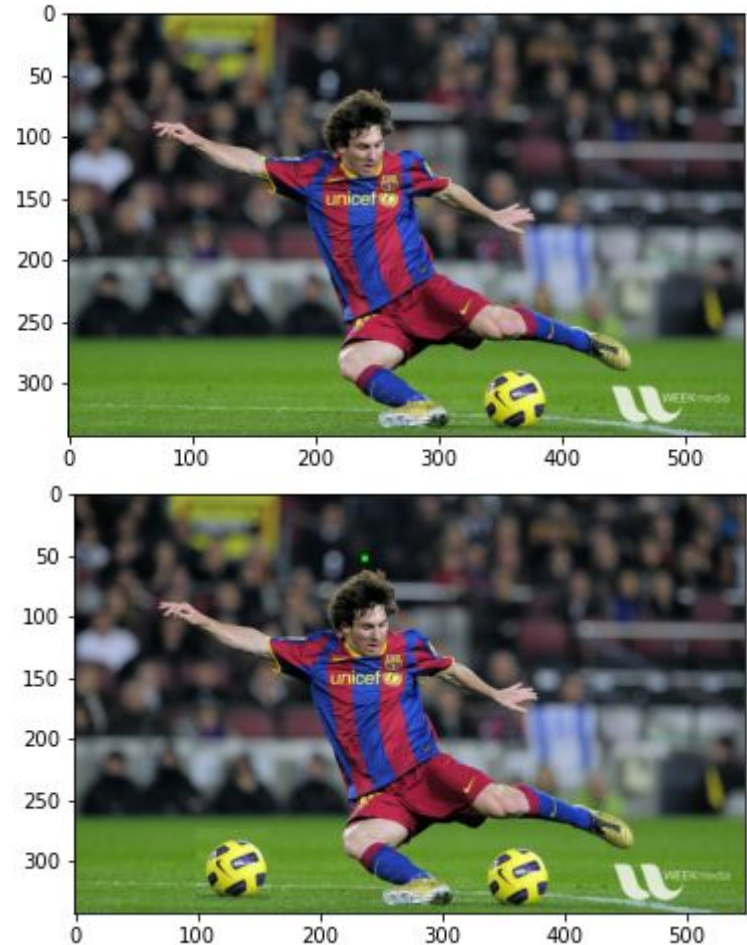
- Import numpy, opencv, and matplotlib.pyplot
- Print their versions
- Using cv2.imread, read one of the images
- Make sure the image is opened by using (is None)
- Print the image type, the shape, and a small portion of the image
- Using matplotlib.pyplot.imshow, show the image

Operations on images

- Create a black image and try out all of these commands
 - `cv2.line(image, startPoint, endPoint, rgb, thickness)`
 - `cv2.rectangle(image, topLeft, bottomRight, rgb, thickness)`
 - `cv2.circle(image, center, radius, rgb, thickness)`
 - `cv2.ellipse(image, center, axes, angle, startAngle, endAngle, rgb, thickness)`
 - `cv2.polylines(image, points, isClosed, rgb, thickness, lineType, shift)`
 - `cv2.putText(image, text, bottomLeft, fontType, fontScale, rgb, thickness, lineType)`
- Show the final image

Modify pixels

- Load an image
- Change the color of a 5 by 5 square to green
- Copy a 60 by 60 square from one location to another
- Show the changed image




Edge detection - Sobel

- Load an image
- Using `cv2.Sobel(src, ddepth, dx, dy, ksize=3, scale=1.0)`,
 - `src`: input image
 - `ddepth`: output image depth
 - `dx`: order of the derivative x
 - `dy`: order of the derivative y
 - `ksize`: size of the extended Sobel kernel; it must be 1, 3, 5, or 7
 - `scale`: optional scale factor for the computed derivative values
- Apply sobel filter in x direction to the image
- Apply sobel filter in y direction to the image
- Find the intensity
- Show the result (use `cmap='gray'`)

Edge detection - Canny

- Load an image
- Using `cv2.Canny(image, threshold1, threshold2, apertureSize=3, L2gradient=False)`
 - image: 8-bit grayscale input image
 - threshold1/threshold2: thresholds for the hysteresis procedure
 - apertureSize: aperture size for the Sobel() operator
 - L2gradient: A flag. True to use $L2$ -norm of gradients. $L1$ -norm for False
- Apply canny edge detector with a few different thresholds
- Show the result (use `cmap='gray'`)



Feature detection

- Read an image
- Create a SIFT feature detector with `cv2.xfeatures2d.SIFT_create()`
- Detect features using `detect(image, section)` on your detector
 - If its named `st`, it would be: `st.detect()`
 - 2nd pos argument is a mask indicating a part of image to be searched in, you don't need to use it, just put `None`
- Draw the detected keypoints using `cv2.drawKeypoints(image, keyPoints, section)`

Feature detection - Continued

- Extract the SIFT descriptor using `sift.compute(image, keyPoints)`
- You can detect features AND extract descriptor at the same time using `sift.detectAndCompute(image, section)`

Feature matching - Try 1

- Read 2 images
 - 'images/box.png', 'images/box_in_scene.png'
- Detect features and extract descriptors for both
- Create a bruteforce matcher using `cv2.BFMatcher(cv2.NORM_L2)`
- Use `knnMatch(des1, des2, k)` on your brute force matcher to get matches
 - This will find k matches for each keypoint in des1
 - Set k to 2 for now
- We should only keep good matches
 - One way is to do a ratio test: if distance of best match/distance of second best match < threshold => keep the best match
 - Use 0.75 threshold for now
- Use `cv2.drawMatches(img1, kp1, img2, kp2, matches, None)` to draw your matches

Feature matching - Try 2

- Read 2 images
 - 'images/box.png', 'images/box_in_scene.png'
- Detect features and extract descriptors for both
- Normalize the descriptors
- Compute Hellinger Distance for all pairs of descriptors in img1 and img2
 - $\sqrt{(1 - \sqrt{\text{des1} \cdot \text{des2}})}$
 - Try to vectorize this
- Use the ratio test to keep good matches
 - You can match 2 features yourself by using `cv2.DMatch(queryIdx, trainIdx, imgIdx, distance)`
 - For example `match = cv2.DMatch(queryIdx, trainIdx, imgIdx, distance)`
- Use `cv2.drawMatches(img1, kp1, img2, kp2, matches, None)` to draw your matches

Face detection

- Using `cv2.CascadeClassifier()`, create a cascade classifier with the file 'detect/haarcascade_frontalface_alt.xml'
 - https://github.com/opencv/opencv/blob/master/data/haarcascades/haarcascade_frontalface_alt.xml
 - <https://www.cs.cmu.edu/~efros/courses/LBMV07/Papers/viola-cvpr-01.pdf>
 - This file has a cascade of simple face features
- Write a function that takes an image path
 - Convert it to grayscale
 - Use `detectMultiScale` on your cascade classifier to detect faces
 - `scaleFactor`: How much the image size is reduced at each image scale=1.2
 - `minNeighbors`: How many neighbors each candidate rectangle should have to retain it=5
 - `minSize`: Minimum possible object size=(30,30)
 - Draw a rectangle around each face
 - Show the image
- Run you function on a couple of images that have faces